

## Assignment 4: Hierarchical Parallelism

Please read these instructions carefully. Your grade is based upon meeting all requirements stated in this assignment.

### Background and Problem Statement

Distributed computing both opens the possibility of an entire network of computers collaborating simultaneously on a single application while also enabling multiple levels of concurrent parallelism. Using MVAPICH2, an open-source standard MPI implementation, combined with global shared memory threading techniques we will implement a hierarchical parallel version of producer/consumer.

Create an MPI version of your producer/consumer program according to the following guidelines:

- Use a total of 5 MPI processes. One MPI process per compute node.
- Follow the guidelines from previous labs, where applicable, unless otherwise noted in this assignment.
- Your rank 0 MPI process will be your master MPI process. This process will:
  - Perform all input and output functions (reading files etc)
  - Perform process-level workload distribution by sending (groups of) work items to each remote MPI process.
  - Use the global communicator for communication between the master MPI process and other MPI processes
  - Report the cmd, encoded key and decoded key values for each transform as with previous labs.
  - Additionally, report the MPI rank and thread id of the producer which performed the encoding, as well as the
  - MPI rank and thread id of the consumer which performed the decoding,
  - Collect and report appropriate elapsed wall-clock timing information (time(1)) for comparison with your other producer/consumer implementations.
- Ranks 1 - 4 will perform parallel transform computations
  - Using OpenMP, create (request) one thread for each node core
  - Design and implement an appropriate thread-level workload distribution using some number (your choice) of
  - producer nodes and some number (your choice) of consumer nodes.
  - Communicate required output values to the master MPI process using the global communicator
  - You may use the global communicator or a custom communicator(s) for communication between transform processes, if needed.

### Program Requirements

- Start with the skeleton program provided and then re-architect it as needed.
- Your final source code MUST be named `prod_cons_mpi_hybrid.c`
  - What to do if vim gives you errors about the swap file being present?
  - When your vim crashes without closing things properly, it leaves behind a hidden .swp file. In your case, it is probably `named. prod_cons_mpi_hybrid.c.swp`.
  - Please note `the dot` in front of the file name. If you delete that swp file, you will not get the error.
- Please find the optimal value of producer and consumer OMP threads that yields the best performance.
- You may read the input from either a serial or a parallel region within your program. All input should be

read from stdin as in OpenMP lab.

- You program should support having multiple “producer” threads which will perform populate a shared data structure like a stack or a queue.
  - The simplest way to implement this would be with a stack with one pointer as below.
  - But you are free to choose any other data structure you prefer. The data structure can be either declared globally or locally to the main function.
    - `int buffer[MAX_BUF_SIZE] = {0};`
    - `int location = 0;`
- You program should support having multiple “consumer” threads, which will take work from the shared data structure, perform the dummy computations.
  - The consumers MUST consume ALL the data the producers put in the shared data structure. Nothing should be left in the shared data structure once all the consumers have exited. If consumers do not read all the data, you will lose points.
  - We use a usleep to simulate this work.
    - The skeleton file has two values for usleep.
    - One is 10 and another is 100000.
    - Use 10 for development.
    - You must use a value of 100000 when doing the final runs for your submission. If you reduce this because your program took too long to finish/did not finish, it will be considered as a sub-optimal implementation, and you will lose points
- You may perform your output from either a serial or a parallel region within your program.
  - Output should be printed on stdout.
  - The order of output NEED NOT be the same as the order of the input file.
- The skeleton program has some prints for output.
- Make sure the final files you submit has the best values for the number OpenMP threads hardcoded in the code.
  - The grader WILL NOT try out your program with different number of OpenMP threads.
  - He will run your program and use the time it gives when he runs it to grade your assignment and give you points.
- The top 10 fastest programs will be subjected to a stress benchmark.
  - The top 5 from that pool will receive extra credits.
- You should try to minimize the number of changes needed.
  - We will be using the “wc -l” command to count the number of lines in each program.

### Input Files Provided

The following file have been provided to you. All files are available on Owens in the `/fs/ess/PAS2171/ /Lab4-Assignment` directory.

- `prod_consumer_orig.c` – The skeleton file
- `run_script.sh` – A script to run the final program and collate the output
- Input files – `longlist` and `shortlist`

## Testing Instructions

- Use the MPI C compiler ("mpicc") only obtained after loading the default ICC-based `mvapich2/2.3.3` module. You should see the following modules available by default `intel/19.0.5` `mvapich2/2.3.3`
  - `run_script.sh` should already do this.
  - Your program should compile without warnings.
  - Points will be deducted for any warnings
- Logging in to Owens using SSH
 

```
ssh -X -C <username>@owens.osc.edu
```
- When compiling and making short (less than a couple minutes) test runs, you can use the login nodes.
- Note that the Owens HPC system only has a limited number of compute cores on each node. If you exceed this, you will see very poor performance as different threads will be competing for CPU resources. Please keep this in mind while choosing the number of OpenMP threads.
- For final benchmark runs or when running with the larger value of `usleep`, please do it on an interactive allocation as described below.

- Getting an interactive allocation

```
salloc -N <number_nodes> -A PAS2171 -p <partition> --time=<time_in_minutes>
```

e.g. To request for one full compute node from the serial partition for 30 minutes, you can use the following command. I am just using 30 as an example. Please use higher values as appropriate.

```
-bash-4.2$ salloc -N 5 -n 140 -A PAS2171 -p debug --time=30
```

```
salloc: Pending job allocation 17868816
```

```
salloc: job 17868816 queued and waiting for resources
```

```
salloc: job 17868816 has been allocated resources
```

```
salloc: Granted job allocation 17868816
```

```
salloc: Waiting for resource configuration
```

```
salloc: Nodes o0646 are ready for job <- That is the node which was allocated for you
```

```
bash-4.2$ hostname
```

```
owens-login03.hpc.osc.edu <- You're still on the head node when you get the control back
```

```
bash-4.2$ ssh o0646 <- This will move you to the node that was allocated
```

```
<....snip....>
```

```
-bash-4.2$ hostname Now you're on the node that was allocated.
```

```
o0646.ten.osc.edu
```

## Report Requirements

Provide a 1–2-page report which includes the following (note that accuracy and brevity are prized, needlessly lengthy reports may lower your score).

- You will submit the PDF report on Carmen.
- In your PDF report, you should address the following points
  - Provide a brief explanation describing how your changes ensure use of shared memory.
  - Provide a chart or graph which illustrates the scalability (or lack thereof) of your solution.

- Y-axis of the graph will represent time taken for execution of the program as reported by the run\_script.sh
- X-axis of the graph will represent the time taken for each number of threads.

– Please choose one value for OMP threads which gave you the best performance.

– Hard code this in your solution so that the grader does not have to pass any additional arguments

– Explain why the values you chose gave the best performance

– Describe any unusual or unexpected results you found.

## Code Submission Instructions

- The run\_script.sh file will compile your program and run it for different combinations of producers and consumers with different input files.
- On successful completion, the run\_script.sh file will create a sub-directory named "cse5441-mpi-lab" and place all necessary files for submission there.
- Submit your cse5441-mpi-lab directory from an Owens login node using the OSC submit system
  - More details about how to submit your assignment from OSC is available at the following website: [https://www.osc.edu/resources/getting\\_started/howto/howto\\_submit\\_homework\\_to\\_repository\\_at\\_osc](https://www.osc.edu/resources/getting_started/howto/howto_submit_homework_to_repository_at_osc)
- Submit this directory to "lab4"
  - You can use the following command to do the submission
    - "/fs/ess/PAS2171/CSE-5441-SP22/submit"
  - Note that the system will automatically close at the specified deadline for submission. You will not be able to submit your files after that.
  - Here is a sample set of steps to submit.

./run\_script.sh

Program may be using....

<...snip...>

```
-rwxr-xr-x 1 osc0577 PZS0622 155864 Apr 14 19:33 a.out
drwxr-xr-x 2 osc0577 PZS0622 4096 Apr 14 19:35 cse5441-mpi-lab
-rw-r--r-- 1 osc0577 PZS0622 6393 Apr 14 16:40 longlist
-rw-r--r-- 1 osc0577 PZS0622 11165 Apr 14 19:22 prod_cons_mpi_hybrid.c
-rw-r--r-- 1 osc0577 PZS0622 8051 Apr 14 21:44 prod_consumer_orig.c
-rwxr-xr-x 1 osc0577 PZS0622 4476 Apr 14 19:33 run_script.sh
-rw-r--r-- 1 osc0577 PZS0622 492 Apr 14 16:40 shortlist
```

-bash-4.2\$ /fs/ess/PAS2171/CSE-5441-SP22/submit

\*\*\*\*\*

Hello. This script will submit your assignment assigned by Hari Subramoni to OSC

Note:

Before you run this script, please create one directory which includes all the files you want to submit

The previous submission of the same assignment from you will be replaced by the new submission

Enter the name of assignment and press [ENTER]

>>>lab4

Enter the absolute path of the directory which includes all the files of your assignment lab2 and press [ENTER]

You will submit the current directory if you simply press [ENTER]

>>>/users/PZS0622/osc0577/CSE-5441/MPI/Homework/cse5441-mpi-lab

Your assignment lab4 is submitted successfully. Thank you!

## Rubrics

Note to students: I have attempted (to the best of my abilities) to cover all possible cases and create a comprehensive rubric. I am and have been updating these based on feedback/questions I have received over the duration of this course. Even after this, it could still not cover all the questions you may have.

The student should have submitted one program – prod\_cons\_mpi\_hybrid.c.

1. **70 points** for the correct program without any data validation errors
  - a. For a correct program, the points will be divided as follows
    - i. 50 points for the correctness of the implementation
    - ii. 20 points for the quality of the implementation
  - b. If you have validation errors, you will lose a minimum of 50 points.
  - c. If the program does not run for the grader, a minimum of 50 points will be deducted.
    - i. Questions of the type “the program ran fine for me, but not for the grader”, “the program did not hang for me, but it hung for the grader” etc may be considered on a case-by-case basis.
  - d. If you forget to submit your program because you did not follow the instructions listed above, you will lose all 70 points that are meant for the program. No excuses.
2. **15 points** for getting approximately linear scaling/improvements in performance with the input longlist on the Owens batch partition.
  - a. The program should approximately show linear scaling/improvements in compute time as reported by the run\_script.sh.
  - b. If the improvement is less than 80% of linear scaling, the student will lose 15 points for the program.
3. **10 points** for the report and the clarity with which the student has described the changes
4. **5 points** if the program cleans up all resources it allocates
  - a. The program should cleanup all resources it allocated.
  - b. If it does not – i.e., if there are memory leaks, then the student will lose points based on the percentage of resources they failed to clear up.
5. **Late submissions will lose points**
  - a. The due date will not be extended

- b. 5 points will be deducted for every 24 hours of delay for a maximum allowable delay of 48 hours.
- c. Beyond 48 hours of delay, the student will receive no points for the assignment.

**Academic Integrity**

Collaborating or completing the assignment with help from others or as a group is NOT permitted

Copying or reusing previous work done by others is not permitted. You may reuse the work you did as part of this class